

Web Usage Mining: Recommendation

Outline

- ▶ Introduction to recommendation
- ▶ Content-based recommendation
- ▶ Collaborative filtering
- ▶ Remarks and practical tips

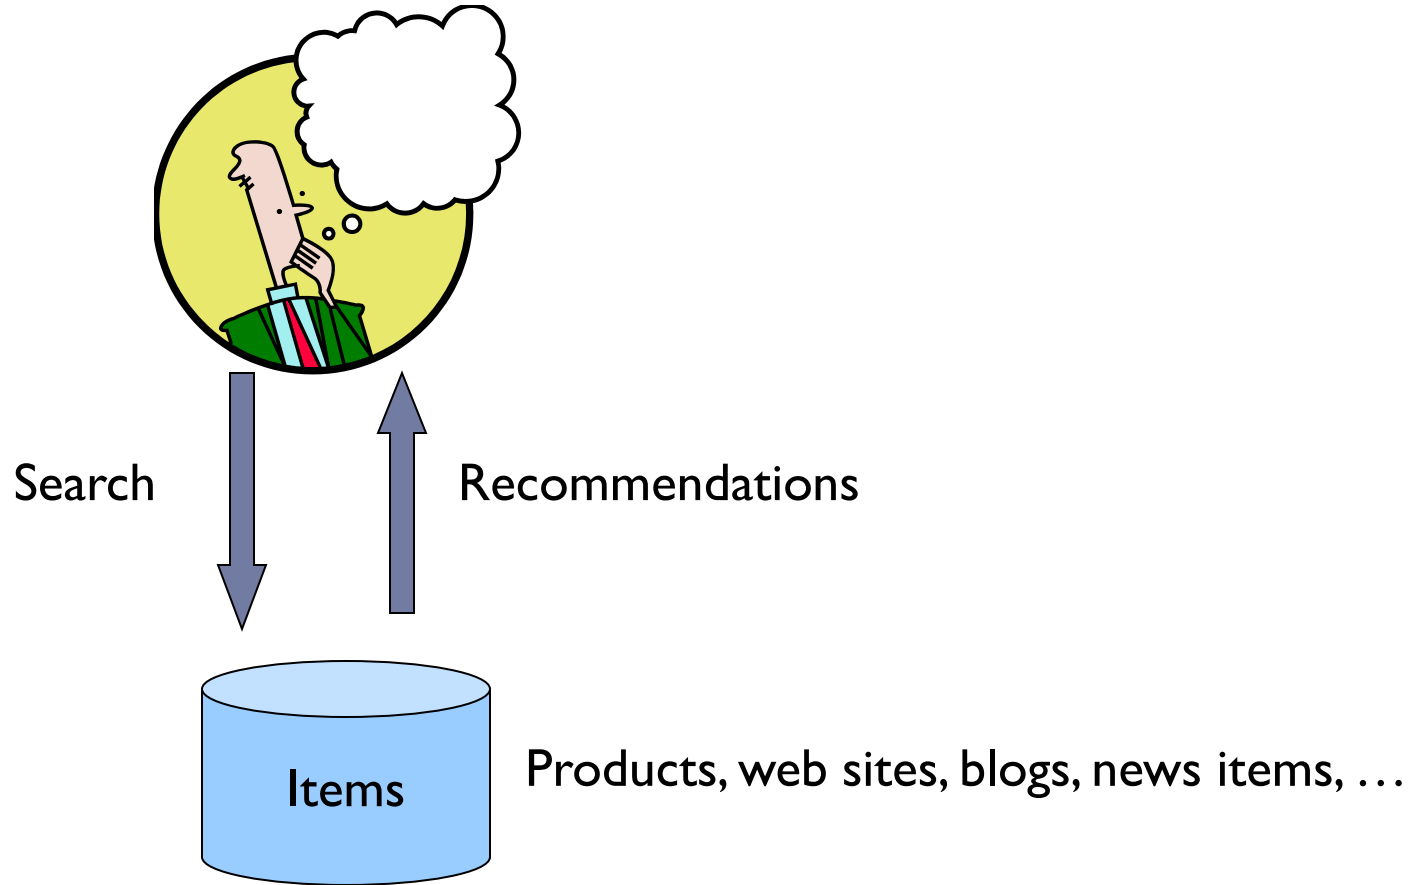
[illegible]

Why Recommendation?

- ▶ Large quantity
- ▶ Diversified quality



Recommendations



From scarcity to abundance

- ▶ Shelf space is a scarce commodity for traditional retailers
 - ▶ Also: TV networks, movie theaters,...
- ▶ The web enables near-zero-cost dissemination of information about products
 - ▶ From scarcity to abundance
- ▶ More choice necessitates better filters
 - ▶ Recommendation engines
 - ▶ How **Into Thin Air** made **Touching the Void** a bestseller (<http://www.wired.com/wired/archive/12.10/tail.html>)



Recommendation Types

- ▶ **Editorial**

- ▶ List of favorites
- ▶ Lists of “essential” items

- ▶ **Simple aggregates**

- ▶ Top 10, Most Popular, Recent Uploads

- ▶ **Tailored to individual users**

- ▶ Amazon, Netflix, Taobao.com...



Formal Model

- ▶ X = set of Customers
- ▶ S = set of Items
- ▶ Utility function $u: X \times S \rightarrow R$
 - ▶ R = set of ratings
 - ▶ R is a totally ordered set
 - ▶ e.g., 0-5 stars, real number in $[0, 1]$



Utility Matrix

	King Kong	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4



Key Problems

- ▶ Gathering “known” ratings for matrix
 - ▶ How to collect the data in the utility matrix
- ▶ Extrapolate unknown ratings from known ratings
 - ▶ Mainly interested in high unknown ratings
 - ▶ We are not interested in knowing what you don't like but what you like
- ▶ Evaluating extrapolation methods
 - ▶ How to measure success/performance of recommendation methods



Gathering Ratings

- ▶ **Explicit**

- ▶ Ask people to rate items
- ▶ Doesn't work well in practice – people can't be bothered

- ▶ **Implicit**

- ▶ Learn ratings from user actions
e.g., purchase implies high rating
- ▶ What about low ratings?



Extrapolating Utilities

- ▶ Key problem: matrix U is **sparse**
 - ▶ Most people have not rated most items
 - ▶ Cold start:
 - ▶ New items have no ratings
 - ▶ New users have no history
- ▶ Three approaches to recommender systems
 - ▶ **Content-based**
 - ▶ **Collaborative**
 - ▶ Latent factor based



Outline

- ▶ Introduction to recommendation
- ▶ Content-based recommendation
- ▶ Collaborative filtering
- ▶ Remarks and practical tips

Content-based Recommendations

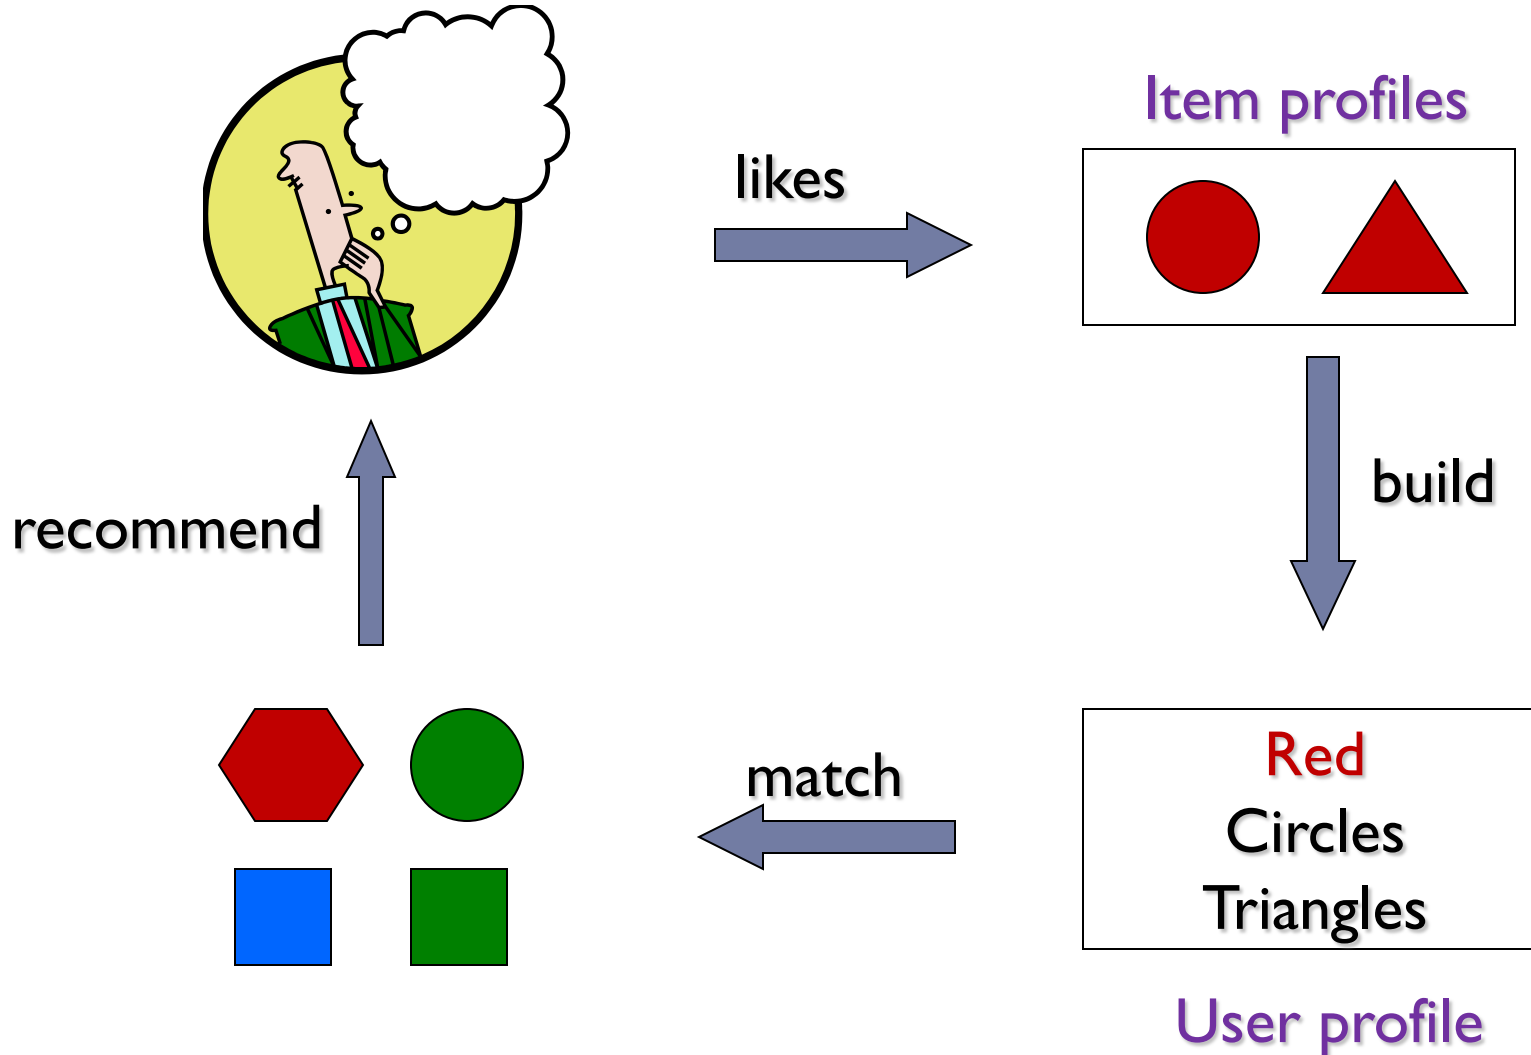
- ▶ Main idea: recommend items to customer x similar to previous items rated highly by x

Example:

- ▶ Movie recommendations
 - ▶ Recommend movies with same actor(s), director, genre, ...
- ▶ Websites, blogs, news
 - ▶ Recommend other sites with “similar” content



Plan of Action



Item Profiles

- ▶ For each item, create an **item profile**
- ▶ Profile is a set (vector) of features
 - ▶ Movies: author, title, actor, director,...
 - ▶ Text: set of “important” words in document
- ▶ How to pick important words?
 - ▶ Usual heuristic from text mining is TF.IDF (Term Frequency * Inverse Doc Frequency)
 - ▶ Term ... Feature
 - ▶ Document ... Item



User Profiles and Prediction

- ▶ User profile possibilities:

- ▶ Weighted average of rated item profiles
- ▶ **Variation:** weight by difference from average rating for item
- ▶ ...

- ▶ Prediction heuristic

- ▶ Given user profile x and item profile i , estimate

$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{\|x\| \cdot \|i\|}$$



Pros: Content-based Approach

- ▶ No need for data on other users
 - ▶ No cold-start or sparsity problems
- ▶ Able to recommend users with unique tastes
- ▶ Able to recommend new and unpopular items
 - ▶ No first-rater problem
- ▶ Able to provide explanations
 - ▶ Can provide explanations of recommended items by listing content-features that caused an item to be recommended



Cons: Content-based Approach

- ▶ Finding the appropriate features is hard
 - ▶ E.g., images, movies, music
- ▶ Overspecialization
 - ▶ Never recommends items outside user's content profile
 - ▶ People might have multiple interests
 - ▶ Unable to exploit quality judgments of other users
- ▶ Recommendations for new users
 - ▶ How to build a user profile?

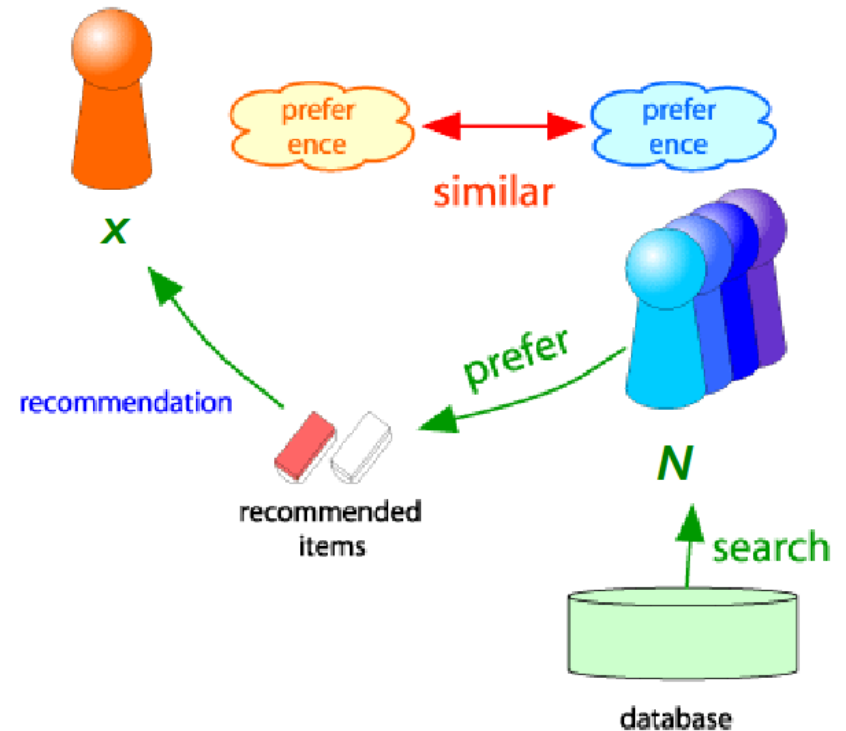


Outline

- ▶ Introduction to recommendation
- ▶ Content-based recommendation
- ▶ Collaborative filtering
- ▶ Remarks and practical tips

Collaborative Filtering (协同过滤)

- ▶ Consider user x
- ▶ Find set N of other users whose ratings are “similar” to x ’s ratings
- ▶ Estimate x ’s ratings based on ratings of users in N



Similar Users

- ▶ Let r_x be the vector of user x 's ratings
- ▶ Jaccard similarity measure
 - ▶ Problem: ignores the value of the rating
- ▶ Cosine similarity measure
 - ▶ $\text{sim}(x,y) = \cos(r_x, r_y)$
 - ▶ Problem: treats missing ratings as “negative”
- ▶ Pearson correlation coefficient
 - ▶ S_{xy} = items rated by both users x and y

$$\text{sim}(x, y) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2 (r_{ys} - \bar{r}_y)^2}}$$



Rating Predictions

- ▶ Let r_x be the vector of user x 's ratings
- ▶ Let N be the set of k users most similar to x who have rated item i
- ▶ Prediction for item s of user x :
 - ▶ $r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$
 - ▶ $r_{xi} = \frac{\sum_{y \in N} s_{xy} r_{yi}}{\sum_{y \in N} s_{xy}}$ $s_{xy} = \text{sim}(x, y)$
 - ▶ Other options?
- ▶ Many other tricks possible...



Item-Item Collaborative Filtering

- ▶ So far: User-user collaborative filtering
- ▶ Another view: Item-item
 - ▶ For item i , find other similar items
 - ▶ Estimate rating for item i based on ratings for similar items
 - ▶ Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

s_{ij} ... similarity of items i and j

r_{xj} ... rating of user u on item j

$N(i;x)$... set of items rated by x similar to i



CF: Common Practice

- ▶ Define similarity s_{ij} of item i and j
- ▶ Select k nearest neighbors $N(i; x)$
 - ▶ Items most similar to i , that were rated by x
- ▶ Estimate rating as the weighted average:

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$

- ▶ b_{xi} : baseline estimate for r_{xi}
- ▶ $b_{xi} = \mu + b_x + b_i$
- ▶ μ =overall mean movie rating
- ▶ b_x =rating deviation of user x = (average rating of user x)- μ
- ▶ b_i =rating deviation of movie i



Item-Item vs. User-User

- ▶ In practice, it has been observed that item-item often works better than user-user
- ▶ Why?
 - ▶ Items are simpler, users have multiple tastes



Pros and Cons of Collaborative Filtering

😊 Works for any kind of item

- ▶ No feature selection needed

😞 Cold start:

- ▶ Need enough users in the system to find a match

😞 Sparsity:

- ▶ The user/ratings matrix is sparse
- ▶ Hard to find users that have rated the same items

😞 First rater:

- ▶ Cannot recommend an item that has not been previously rated
- ▶ New items, esoteric items

😞 Popularity bias:

- ▶ Cannot recommend items to someone with unique taste
- ▶ Tends to recommend popular items



Hybrid Methods

- ▶ Implement or more different recommenders and combine predictions
- ▶ Add content-based methods to collaborative filtering
 - ▶ Item profiles for new item problem
 - ▶ Demographics (人口统计学) to deal with new user problem



Outline

- ▶ Introduction to recommendation
- ▶ Content-based recommendation
- ▶ Collaborative filtering
- ▶ Remarks and practical tips

Evaluation

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					



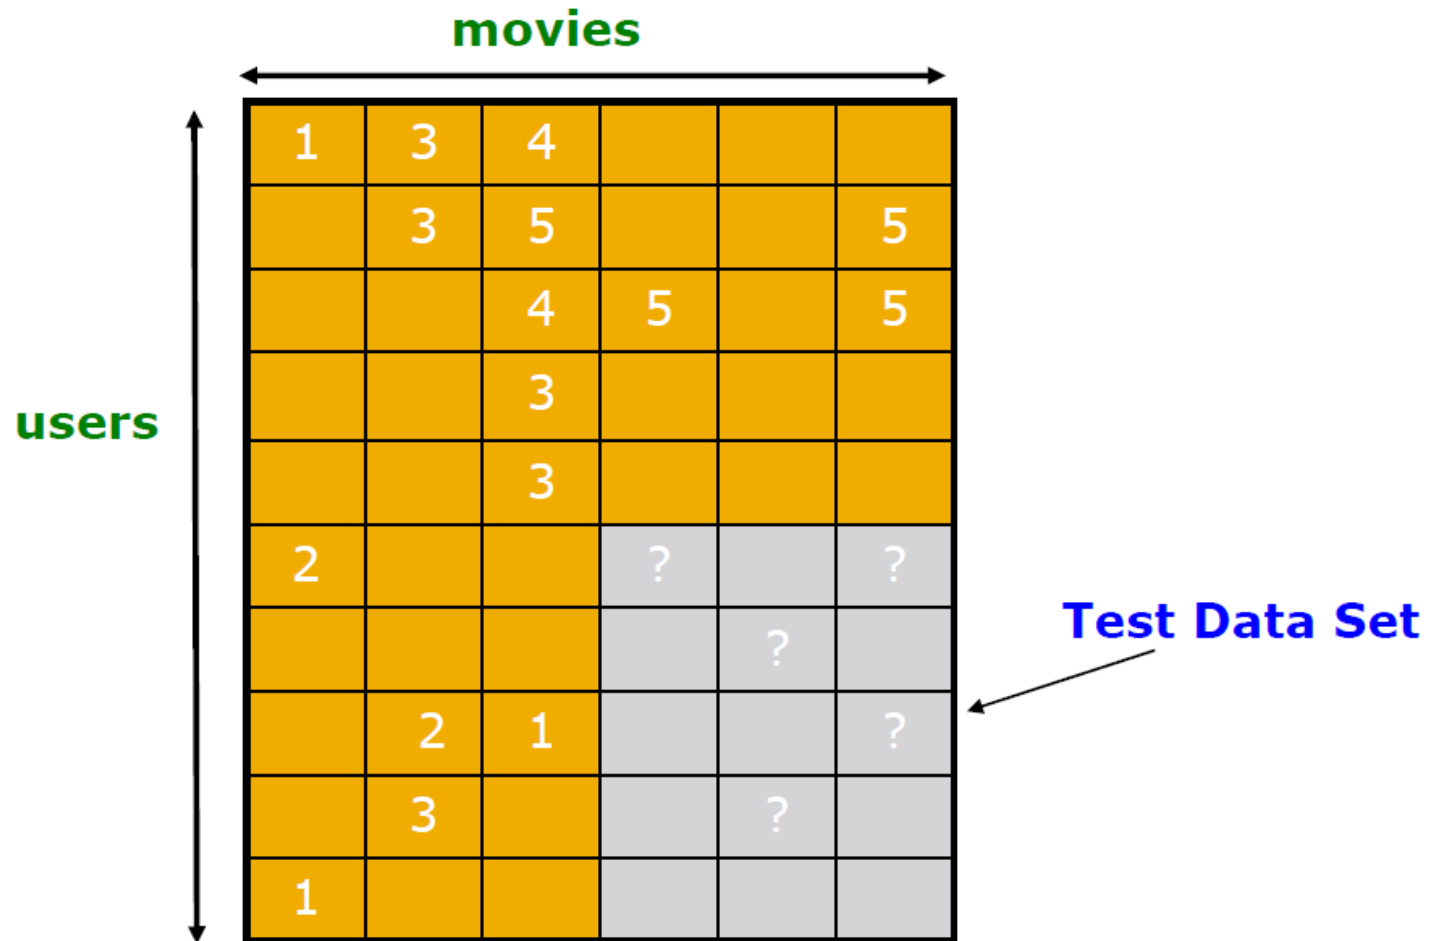
Evaluation

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

Test Data Set



The matrix shows ratings for 10 users across 6 movies. The top 5 rows and first 3 columns are orange, while the bottom 5 rows and last 3 columns are gray, labeled as the 'Test Data Set'. Ratings are integers from 1 to 5, or '?' for unknown values.

Evaluating Predictions

- ▶ Compare predictions with known ratings
 - ▶ Root-mean-square error (RMSE)
 - ▶ $\sqrt{\sum_{xi}(r_{xi} - r_{xi}^*)^2}$ where r_{xi} is predicted, r_{xi}^* is the true rating
 - ▶ Precision at top 10:
 - ▶ % of those in top 10
- ▶ Another approach: 0/1 model
 - ▶ Coverage
 - ▶ Number of items/users for which system can make predictions
 - ▶ Precision
 - ▶ Accuracy of predictions
 - ▶ Receiver operating characteristic (ROC)
 - ▶ Tradeoff curve between false positives and false negatives



Problems with Measures

- ▶ **Narrow focus on accuracy sometimes misses the point**
 - ▶ Prediction Diversity
 - ▶ Prediction Context (情境)
 - ▶ Order of predictions
- ▶ **In practice, we care only to predict high ratings**
 - ▶ RMSE might penalize a method that does well for high ratings and badly for others



Collaborative Filtering: Complexity

- ▶ Expensive step is finding k most similar customers:
 $O(|X|)$
- ▶ Too expensive to do at runtime
 - ▶ Need to pre-compute
- ▶ Naïve pre-computation takes time $O(|S| \cdot |X|)$
 - ▶ Near-neighbor search in high dimensions
- ▶ Can use clustering, partitioning as alternatives, but quality degrades



The Netflix Prize

- ▶ **Training data**

- ▶ 100 million ratings, 480,000 users, 17,770 movies
- ▶ 6 years of data: 2000-2005

- ▶ **Test data**

- ▶ Last few ratings of each user (2.8 million)
- ▶ Evaluation criterion: root mean squared error (RMSE)
- ▶ Netflix Cinematch RMSE: 0.9514

- ▶ **Competition**

- ▶ 2700+ teams
- ▶ \$1 million prize for 10% improvement on Cinematch



Major Challenges in Recommendation

- ▶ Data Sparsity
- ▶ Scalability
- ▶ Cold Start
- ▶ Diversity vs. Accuracy
- ▶ Vulnerability to Attacks
- ▶ Value of Time
- ▶ Evaluation of Recommendations
- ▶ User Behavior mining
- ▶ User Interface
- ▶ Social-based recommendation
- ▶ Multi-resource data



Factors that Influence Recommendation

- ▶ Temporal
- ▶ Spatial / location
- ▶ Social
- ▶ Trust

...



Some Works of Web Mining

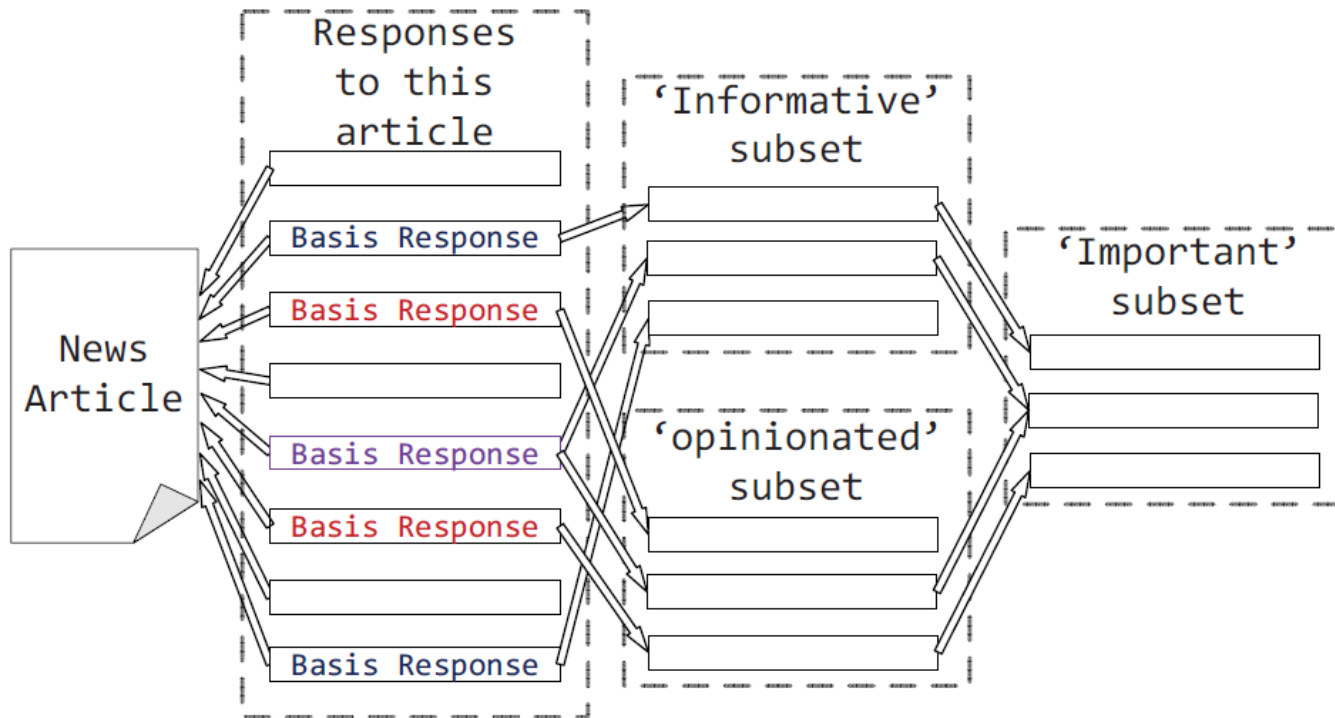
Web Mining Results

- Map vocabularies of **two languages** into a shared vector space

Chinese-English word embeddings

Web Mining Results

► Selecting social media responses to news



Web Mining Results

► Video description with bullet-screen comments

				
22:11	23:06	24:20	25:27	25:51
Nian Nian has slanted eyebrows.	Mr. Su used to be Brother Su. How sad.	I feel sorry for both Jing Rui and Su from their conversation.	Mr. Su lost a good friend forever!	Jing Rui leaves the place where his heart broke and dreams faded away.
keywords by LDA: 'Su', 'say', 'really', 'Jing Rui', 'leave', 'eyebrow', 'Nian Nian', 'come', 'friends', 'love'				

Web Mining Results

► Community detection



A word cloud of biological terms. The words are arranged in a roughly circular shape. The most prominent words are 'islet', 'mice', 'antibodi', 'beta', 'tcell', 'hla', 'suscept', 't', 'nod', 't1d', 'cd4', 'autoantibodi', 'antigen', 'autoimmun', and 'cell'. The words are colored in blue, red, and black.

islet
mice
antibodi
lymphocyt
beta
hla
tcell
suscept
t
nod
t1d
cd4
autoantibodi
antigen
autoimmun
cell



A word cloud of biological terms. The words are arranged in a roughly circular shape. The most prominent words are 'autoimmun', 'cell', 'mice', 'gene', 'express', 'betacel', 't', 'beta', 't1d', 'rat lymphocyt', 'patient', 'suscept', 'antigen', 'iddm', 'diseas', 'antibodi', 'nod', and 'islet'. The words are colored in blue, red, and black.

autoimmun
cell
mice
gene
express
betacel
t
beta
t1d
rat lymphocyt
patient
suscept
antigen
iddm
diseas
antibodi
nod
islet

Web Mining Results

► Recommendation across multiple domains

Domains	Book	Movie
#Users	13090	13090
#Items	17590	17922
Sparsity	99.66%	98.68%

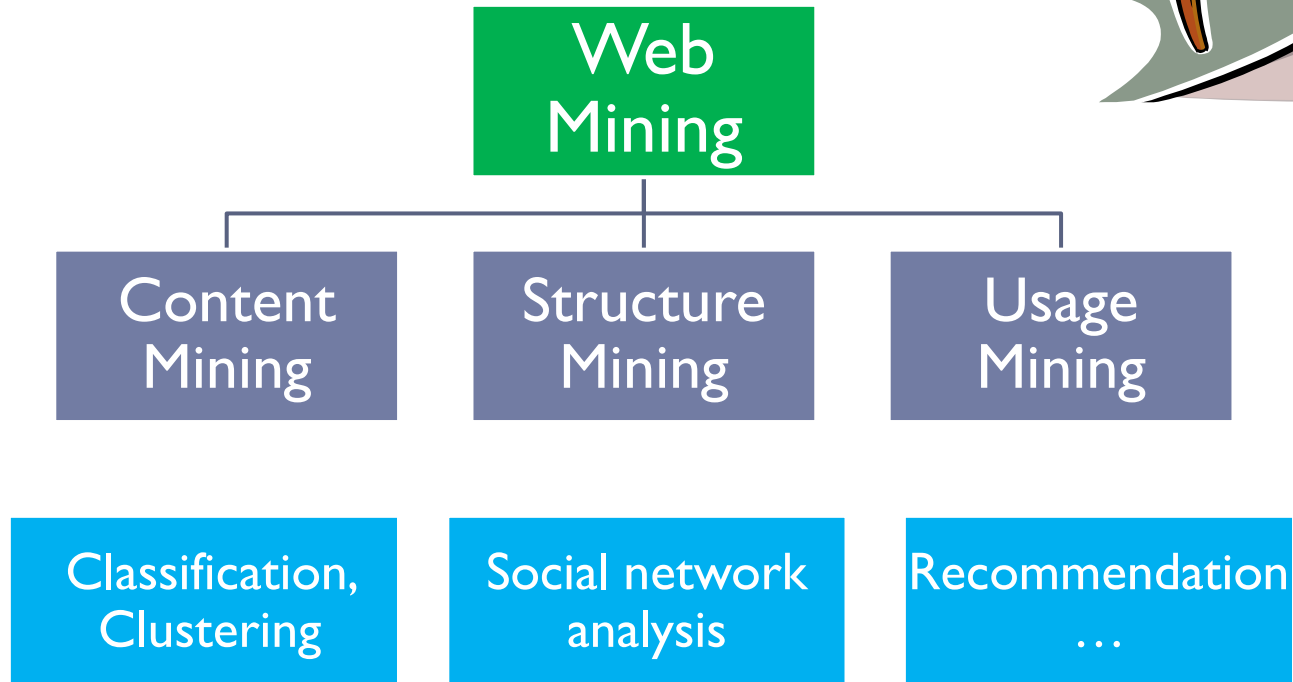
Domains	Training	SVP	TNNR-PG	PMF	CMF	GSMF	Aligned MC
Book	80%	0.9606(0.4898)	0.8801 (0.6144)	0.7809 (0.5235)	0.8172 (0.6362)	0.7813 (0.5684)	0.7389 (0.4008)
	60%	1.0147(0.4658)	0.9066 (0.5663)	0.7967 (0.5353)	0.8517 (0.6523)	0.7962(0.6078)	0.7479 (0.4550)
	40%	1.1571(0.4175)	1.0239 (0.5563)	0.8397 (0.5083)	0.9345 (0.6227)	0.8030 (0.5643)	0.7558 (0.4911)
Movie	80%	0.7661(0.6011)	0.7336 (0.6524)	0.7342 (0.6014)	0.7325 (0.6228)	0.7315 (0.6177)	0.7130 (0.6367)
	60%	0.7870(0.5905)	0.7429 (0.6391)	0.7432 (0.5952)	0.7423 (0.6142)	0.7401(0.5978)	0.7209 (0.6643)
	40%	0.8387(0.5616)	0.7752 (0.6259)	0.7678 (0.5764)	0.7829 (0.5784)	0.7870 (0.4892)	0.7342 (0.6885)





Summary of Web Mining

Roadmap



Note: Helpful to combine usage with content and structure

Introduction

- ▶ 网络挖掘的概念，包含哪些方面的内容，分别有哪些重要应用？

Data

- ▶ 概念：数据对象(Objects)，属性(Attributes)，维度(Dimensions)，特征(features)
- ▶ 高维诅咒(Curse of dimensionality)现象。
- ▶ 对于数据的预处理有哪些方法？其中需要掌握采样(Sampling)，特征选择(Feature selection)及降维(Dimensionality reduction)的基本原理。



Classification

- ▶ 监督学习(Supervised learning)与无监督学习(Unsupervised learning)的关系与区别。
- ▶ 分类(Classification)的基本原理。
- ▶ 数据的向量表示(Vector space representation)
- ▶ 熟练掌握k近邻算法，包括影响算法性能的要素——近邻个数及距离（相似度）度量。
- ▶ 熟练掌握Logistic regression分类方法。
- ▶ 如何评价分类效果？理解训练错误率，测试错误率以及泛化错误率的区别。



Clustering

- ▶ 聚类(Clustering)的基本原理及准则。
 - ▶ High similarity within clusters
 - ▶ Low similarity between clusters
 - ▶ Important issues of clustering:
 - ▶ Number of clusters, Similarity measure
- ▶ 层次式聚类算法流程，两个类之间的距离定义。
- ▶ 熟练掌握K-means算法——算法流程，优化目标，收敛性分析。
- ▶ 聚类算法的评价标准。
 - ▶ With or without ground truth



Networks: Community

- ▶ 社区(Community)的概念
- ▶ 社区发现与聚类的关系。
- ▶ 如何计算结构相似度？
- ▶ 图分析的一些重要矩阵：邻接(Affinity)矩阵，拉普拉斯(Laplacian)矩阵，以及它们的一些重要性质。
- ▶ Cut概念；ratio cut以及normalized cut的定义及推导。
- ▶ Modularity概念及其推导。与spectral clustering的相同点及不同点。

Point: analyze the eigenvectors of a matrix to explore the structure of the graph



Networks: Influence

- ▶ 几种度量节点中心性的标准。
- ▶ 两种影响力传播模型——线性阈值模型(Linear Threshold Model), 层级传播模型(Independent Cascade Model)的传播过程及区别。
- ▶ 最大影响节点集(Most influential set)——问题建模, 贪心算法以及算法的近似度。
- ▶ 子模性质(submodularity)。



Recommendation

- ▶ 推荐问题的形式。
- ▶ 基于内容的推荐：主要思想。
- ▶ 协同过滤：主要思想和基本方法。
- ▶ 各自的优缺点。

