

k

Learning



Outline

- 1 Reinforcement Learning
- 2 Markov Decision Process
- 3 Learning Algorithms
- 4 Q-Learning in R
- 5 Wrap-Up

Outline

- 1** Reinforcement Learning
- 2 Markov Decision Process
- 3 Learning Algorithms
- 4 Q-Learning in R
- 5 Wrap-Up

Branches of Machine Learning

Supervised Learning


- ▶ Learns from pairs of input and desired outcome (i. e. **labels**)

Unsupervised Learning

- ▶ Tries to find **hidden** structure in **unlabeled** data

Reinforcement Learning

- ▶ Learning from **interacting** with the **environment**
- ▶ No need for pairs of input and correct outcome
- ▶ Feedback restricted to a **reward** signal
- ▶ **Mimics human-like learning** in actual environments



Example: Backgammon

Reinforcement learning can reach a level similar to the top three human players in backgammon

Learning task

- ▶ Select best move at arbitrary board states
→ i. e. with highest probability to win

Training signal


- ▶ Win or loss of overall game

Training

- ▶ 300,000 games played against the system itself

Algorithm


- ▶ Reinforcement learning (plus neural network)



→ Tesauro (1995): Temporal Difference Learning and TD-Gammon. In: Comm. of the ACM, 38:3, pp. 58–68

Reinforcement Learning


- ▶ An **agent interacts** with its environment
- ▶ Agent takes actions that affect the **state** of the environment
- ▶ Feedback is limited to a **reward** signal that indicates how well the agent is performing
- ▶ Goal: improve the behavior given only this limited feedback



Examples

- ▶ Defeat the world champions at backgammon or Go
- ▶ Manage an investment portfolio
- ▶ Make a humanoid robot walk

Agent and Environment



At each step t , the **agent**:

- ▶ Executes action a_t
- ▶ Receives observation s_t
- ▶ Receives scalar reward r_t

The **environment**:

- ▶ Changes upon action a_t
- ▶ Emits observation s_{t+1}
- ▶ Emits scalar reward r_{t+1}


- ▶ Time step t is **incremented** after each iteration

Agent and Environment

Example

- | | | |
|---|-------------|--|
| ① | ENVIRONMENT | ▶ You are in state 3 with 4 possible actions |
| ② | AGENT | ▶ I'll take action 2 |
| ③ | ENVIRONMENT | ▶ You received a reward of 5 units |
| ⋮ | | ▶ You are in state 1 with 2 possible actions |
| ⋮ | | ⋮ |

Formalization



Reinforcement Learning Problem

Finding an optimal behavior

- ▶ Learn optimal behavior π based on **past actions**
- ▶ **Maximize the expected cumulative reward** over time

Challenges

- ▶ Feedback is **delayed**, not instantaneous
- ▶ Agent must reason about the **long-term consequences** of its actions

Illustration

- ▶ In order to maximize one's future income, one has to study now
- ▶ However, the immediate monetary reward from this might be negative

⇒ How do we learn optimal behavior?

Trial-and-Error Learning

The agent should discover optimal behavior via **trial-and-error learning**

1 Exploration

- ▶ Try new or non-optimal actions to learn their reward
- ▶ Gain a better understanding of the environment

2 Exploitation

- ▶ Use current knowledge
- ▶ This might not be optimal yet, but should deviate only slightly

Examples

1 Restaurant selection

- ▶ **Exploitation:** go to your favorite restaurant
- ▶ **Exploration:** try a new restaurant


2 Game playing

- ▶ **Exploitation:** play the move you believe is best
- ▶ **Exploration:** play an experimental move

ϵ -Greedy Action Selection

Idea

- ▶ Provide a **simple heuristic to choose** between exploitation and exploration
- ▶ Implemented via a **random number** $0 \leq \epsilon \leq 1$
 - ▶ With probability ϵ , try a random action
 - ▶ With probability $1 - \epsilon$, choose the current best



- ▶ Typical choice is e. g. $\epsilon = 0.1$
- ▶ Other variants decrease this value over time
→ i. e. agent gains confidence and thus needs less exploration

Outline

- 1 Reinforcement Learning
- 2 Markov Decision Process**
- 3 Learning Algorithms
- 4 Q-Learning in R
- 5 Wrap-Up

Markov Decision Process

- ▶ A **Markov decision process** (MDP) specifies a setup for reinforcement learning
- ▶ MDPs allow to model decision making in situations where outcomes are partly random and partly under the control of a decision maker

Definition

- 1** A Markov Decision Process is a 4-tuple (S, A, R, T) with
 - ▶ A set of possible world **states** S
 - ▶ A set of possible **actions** A
 - ▶ A real-valued **reward function** R
 - ▶ **Transition probabilities** T
- 2** A MDP must fulfill the so-called **Markov property**
 - ▶ The effects of an action taken in a state **depend only on that state** and not on the prior history

Markov Decision Process

State

- ▶ A state s_t is a representation of the environment at time step t
- ▶ Can be directly observable to the agent or hidden

Actions

- ▶ At each state, the agent is able to perform an action a_t that affects the subsequent state of the environment s_{t+1}
- ▶ Actions can be any decisions which one wants to learn

Transition probabilities

- ▶ Given a current state s , a possible subsequent state s' and an action a
- ▶ The transition probability $T_{ss'}^a$ from s to s' is defined by

$$T_{ss'}^a = P[s_{t+1} = s' \mid s_t = s, a_t = a]$$

Rewards

- ▶ A reward r_{t+1} is a **scalar feedback** signal emitted by the environment
- ▶ Indicates how well agent is performing when reaching step $t + 1$
- ▶ The **expected reward** $R_{ss'}^a$ when moving from state s to s' via action a is given by

$$R_{ss'}^a = E[r_{t+1} \mid s_t = s, a_t = a, s_{t+1} = s']$$


Examples

- 1 Playing backgammon or Go
 - ▶ Zero reward after each move
 - ▶ A positive/negative reward for winning/losing a game
- 2 Managing an investment portfolio
 - ▶ A positive reward for each dollar left in the bank

Goal: maximize the **expected cumulative reward** over time

Markov Decision Process

Example: Moving a pawn to a destination on a grid



→ available actions $A(s)$
depend on current state s

- ▶ States $S = \{s_0, s_1, \dots, s_7\}$
- ▶ Actions $A = \{\text{up, down, left, right}\}$
- ▶ Transition probabilities
 - ▶ $T_{s_0, s_3}^{\text{up}} = 0.9$
 - ▶ $T_{s_0, s_1}^{\text{right}} = 0.1$
 - ▶ ...
- ▶ Rewards
 - ▶ $R_{s_6, s_7}^{\text{right}} = +10$
 - ▶ $R_{s_2, s_4}^{\text{up}} = -10$
 - ▶ Otherwise $R = 0$
- ▶ Start in s_0
- ▶ Game over when reaching s_7

Policy

Learning task of an agent

- ▶ Execute actions in the environment and observe results, i. e. rewards
- ▶ Learn a **policy** $\pi : S \rightarrow A$ that works as a **selection function** of choosing an action given a state
- ▶ A policy fully defines the **behavior of an agent**, i. e. its actions
- ▶ MDP policies depend only on the current state and not its history
- ▶ Policies are **stationary** (i. e. time-independent)

Objective

- ▶ Maximize the **expected cumulative reward** over time
- ▶ The expected cumulative reward from an initial state s with policy π is

$$J_{\pi}(s) = \sum_t R_{s_t, s_{t+1}}^{a_t} = E_{\pi} \left[\sum_t r_t \mid s_0 = s \right]$$

Value Functions

Definition

- ▶ The **state-value function** $V_{\pi}(s)$ of an MDP is the **expected reward** starting from state s , and then following once policy π
- ▶ $V_{\pi}(s) = E_{\pi} [J_{\pi}(s_t) \mid s_t = s]$
- ▶ Quantifies how good is it to be in a particular state s

Definition

- ▶ The **state-action value function** $Q_{\pi}(s,a)$ is the **expected reward** starting from state s , taking action a , and then following policy π
- ▶ $Q_{\pi}(s,a) = E_{\pi} [J_{\pi}(s_t) \mid s_t = s, a_t = a]$
- ▶ Quantifies how good is it to be in a particular state s and apply action a , and afterwards follow policy π

Now, we can formalize the **policy** definition (with **discount factor** γ) via

$$\pi(s) = \arg \max_a \sum_{s'} T_{ss'}^a (R_{ss'}^a + \gamma V_{\pi}(s'))$$

Optimal Value Functions

- ▶ While π can be any policy, π^* denotes the optimal one with the highest expected cumulative reward
- ▶ The optimal value functions specify the best possible policy
- ▶ A MDP is solved when the optimal value functions are known

Definitions

- 1** The **optimal state-value function** $V_{\pi^*}(s)$ maximizes the expected reward over all policies

$$V_{\pi^*}(s) = \max_{\pi} V_{\pi}(s)$$

- 2** The **optimal action-value function** $Q_{\pi^*}(s,a)$ maximizes the action-value function over all policies

$$Q_{\pi^*}(s,a) = \max_{\pi} Q_{\pi}(s,a)$$

Outline

- 1 Reinforcement Learning
- 2 Markov Decision Process
- 3 Learning Algorithms**
- 4 Q-Learning in R
- 5 Wrap-Up

Types of Learning Algorithms

Aim: find optimal policy and value functions

Model-based learning

- ▶ Aim: find optimal policy and value functions
- ▶ Model of the environment is as **MDP with transition probabilities**
- ▶ Approach: **learn the MDP model** or an approximation of it

Model-free learning

- ▶ Explicit model of the environment model is not available
→ i. e. transition probabilities are **unknown**
- ▶ Approach: derive the optimal policy **without explicitly formalizing the model**

Outline

- 3** Learning Algorithms
 - Model-Based Learning
 - Model-Free Learning

Model-Based Learning: Policy Iteration

Approach via policy iteration

- ▶ Given an initial policy π_0
- ▶ **Evaluate policy** π_i to find the corresponding value function V_{π_i}
- ▶ **Improve policy** over V_{π} via greedy exploration
- ▶ Policy iteration always **converges to optimal policy** π^*

Illustration

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \dots \xrightarrow{E} V_{\pi^*} \xrightarrow{I} \pi^*$$

with

- ▶ E : policy evaluation
- ▶ I : policy improvement

Policy Evaluation

- Computes the state-value function V_π for an arbitrary policy π via

$$\begin{aligned} V_\pi(s) &= E_\pi [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s_t = s] \\ &= E_\pi [r_{t+1} + \gamma V_\pi(s_{t+1}) \mid s_t = s] \\ &= \sum_a \pi(s, a) \sum_{s'} T_{ss'}^a [R_{ss'}^a + \gamma V_\pi(s')] \end{aligned}$$

- System of $|S|$ linear equations with $|S|$ unknowns
- Solvable but computational expensive if $|S|$ is large
- Advanced methods are available, e. g. iterative policy evaluation

Discount factor

- If $0 < \gamma < 1$, makes cumulative reward finite
- Necessary for setups with infinite time horizons
- Puts more importance on first learning steps, but less on later ones

Iterative Policy Evaluation

- ▶ Iterative policy evaluation uses dynamic programming
- ▶ Iteratively approximate V_π
- ▶ Choose V_0 arbitrarily
- ▶ Then use Bellman equation as an update rule

$$\begin{aligned} V_{k+1}(s) &= E_\pi [r_{t+1} + \gamma V_k(s+1) \mid s_t = s] \\ &= \sum_a \pi(s,a) \sum_{s'} T_{ss'}^a [R_{ss'}^a + \gamma V_k(s')] \end{aligned}$$

- ▶ Sequence V_k, V_{k+1}, \dots converges to V_π as $k \rightarrow \infty$

Policy Improvement

- ▶ Policy evaluation determines the value function V_π for a policy π
- ▶ The alternative step exploits this knowledge to **select the optimal action** in each state
- ▶ For that, policy improvement **searches policy π' that is as good as or better than π**
- ▶ Remedy is to use state-action value function via

$$\begin{aligned}\pi'(s) &= \arg \max_a Q_\pi(s, a) \\ &= \arg \max_a E[r_{t+1} + \gamma V_k(s+1) \mid s_t = s] \\ &= \arg \max_a \sum_{s'} T_{ss'}^a [R_{ss'}^a + \gamma V_k(s')]\end{aligned}$$

- ▶ Afterwards, **continue** with policy evaluation and policy improvement **until a desired convergence** criterion is reached

Outline

- 3** Learning Algorithms
 - Model-Based Learning
 - Model-Free Learning

Model-Free Learning

Drawbacks of model-based learning

- ▶ Requires MDP, i. e. **explicit model** of the dynamics in the environment
- ▶ Transition probabilities are often not available or difficult to define
- ▶ Model-based learning is thus often intractable even in “simple” cases

Model-free learning

- ▶ Idea: **learn directly from interactions** with the environment
- ▶ Only use experience from the sequences of states, action, and rewards

Common approaches

- 1 Monte Carlo methods** are simple but has **slow convergence**
- 2 Q-learning** is more **efficient** due to off-policy learning

Monte Carlo Method

- ▶ Monte Carlo methods require **no knowledge of transition** as in MDPs
- ▶ Perform reinforcement learning from a **sequence of interactions**
- ▶ Mimic policy iteration to find optimal policy
- ▶ Estimate the **value of each action** $Q(s,a)$ instead of $V(s)$
- ▶ Store **average rewards** in state-action table

Example

- ▶ **State-action table**

State	Actions		Optimal Policy
	a_1	a_2	
s_1	2	1	a_1
s_2	1	3	a_2
s_3	2	4	a_2

Monte Carlo Method

Algorithm

- 1 Start with an arbitrary state-action table (and corresponding policies)
→ Often all rewards are **initially set to zero**
- 2 Observe first state
- 3 Choose an action according to **ϵ -greedy action selection**, i. e.
 - ▶ With probability ϵ , pick a **random action**
 - ▶ Otherwise, take **action with highest expected reward**
- 4 **Update state-action table** with new reward (averaging)
- 5 Observe new state
- 6 Go to step 3

Disadvantage

- ▶ High computational time and thus **slow convergence**
→ Method must frequently evaluate a suboptimal policy

Q-Learning

- ▶ One of the most important breakthroughs in reinforcement learning
- ▶ **Off-policy** learning concept
 - ▶ Explore the environment and **at the same time** exploit the current knowledge
- ▶ In each step, take a look forward to the next state and **observe the maximum possible reward for all available actions** in that state
- ▶ Use this knowledge to update the action-value of the corresponding action in the current state
- ▶ Apply **update rule** with **learning rate** α ($0 < \alpha \leq 1$)

$$Q(s,a) \leftarrow \underbrace{Q(s,a)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \left[\underbrace{r'}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \underbrace{\max_{a'} Q(s',a')}_{\text{expected optimal value}} - \underbrace{Q(s,a)}_{\text{old value}} \right]$$

- ▶ Q-learning is repeated for different **episodes** (e. g. games, trials, etc.)

Q-Learning

Algorithm

- 1 Initialize the table $Q(s,a)$ to zero for all state-action pairs (s, a)
- 2 Observe the current state s
- 3 Repeat until convergence
 - ▶ Select an action a and apply it
 - ▶ Receive immediate reward r
 - ▶ Observe the new state s'
 - ▶ Update the table entry for $Q(s,a)$ according to

$$Q(s,a) \leftarrow Q(s,a) + \alpha \left[r + \gamma \max_{a'} Q(s',a') - Q(s,a) \right]$$

- ▶ Move to next state, i. e. $s \leftarrow s'$

Outline

- 1 Reinforcement Learning
- 2 Markov Decision Process
- 3 Learning Algorithms
- 4 Q-Learning in R
- 5 Wrap-Up**

Wrap-Up

Summary

- ▶ Reinforcement learning learns through **trial-and-error** from interactions
- ▶ The **reward** indicates the performance of the agent
→ But without showing how to improve its behavior
- ▶ Learning is grouped into **model-based** and **model-free** strategies
- ▶ A common and efficient model-free variant is **Q-learning**
- ▶ Similar to **human-like learning** in real-world environments
- ▶ Common for trade-offs between long-term vs. short-term benefits

Drawbacks

- ▶ Can be computational expensive when state-action space is large
- ▶ **No R library** is yet available for model-free learning